

You Don't Know Node

Quick Intro to 5 Core Features

This is Sparta! hands-on presentation



Slides & Code



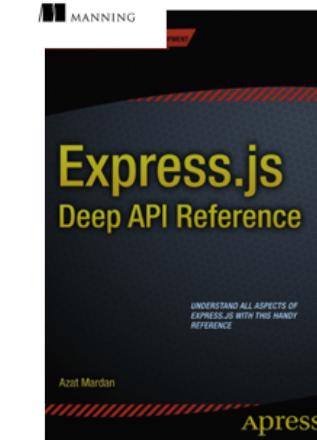
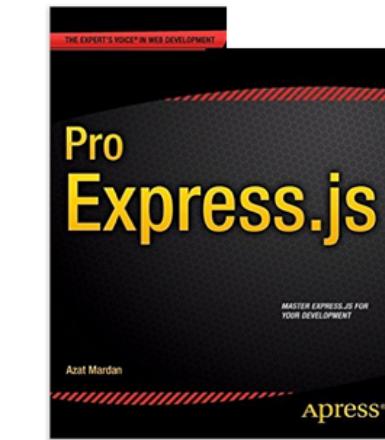
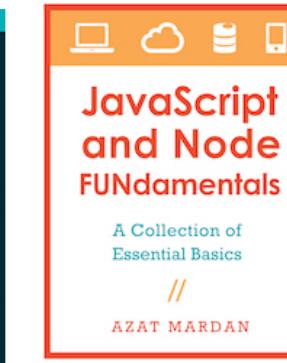
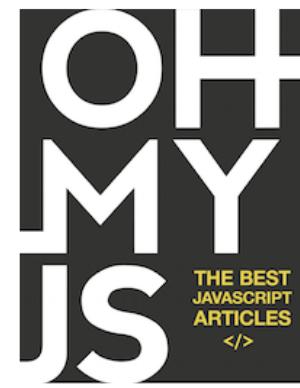
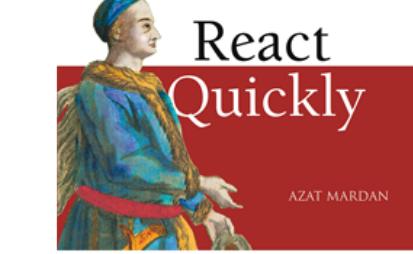
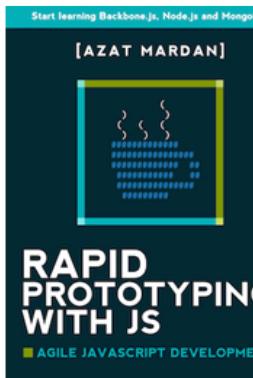
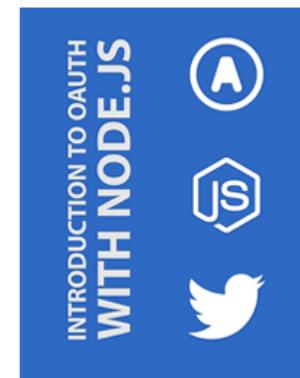
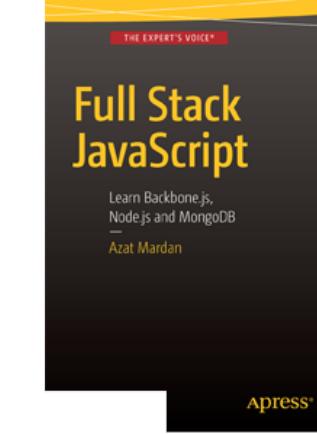
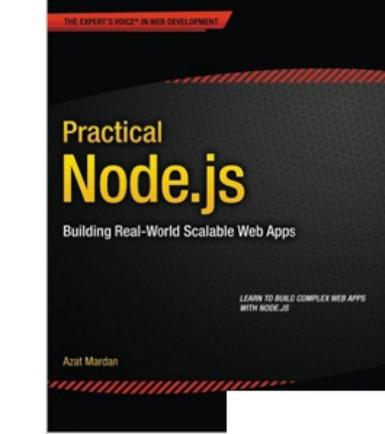
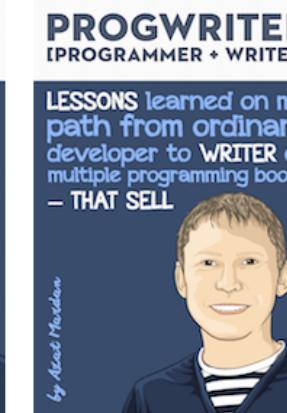
Everything: <https://github.com/azat-co/you-dont-know-node>

or

just PDF: <http://bit.ly/1VJWpQK>

Better Apps—Better Life

About Presenter



Azat Mardan



Twitter: @azat_co
Email: hi@azat.co
Blog: webapplog.com

About Presenter

- » Work: Technology Fellow at Capital One (kind of a big deal)
- » Experience: FDIC, NIH, DocuSign, HackReactor and Storify
- » Books: React Quickly, Practical Node.js, Pro Express.js, Express.js API and 8 others
- » Teach: NodeProgram.com

Capital One in Top 10 US Banks



Starting with basics: Why Use Node?

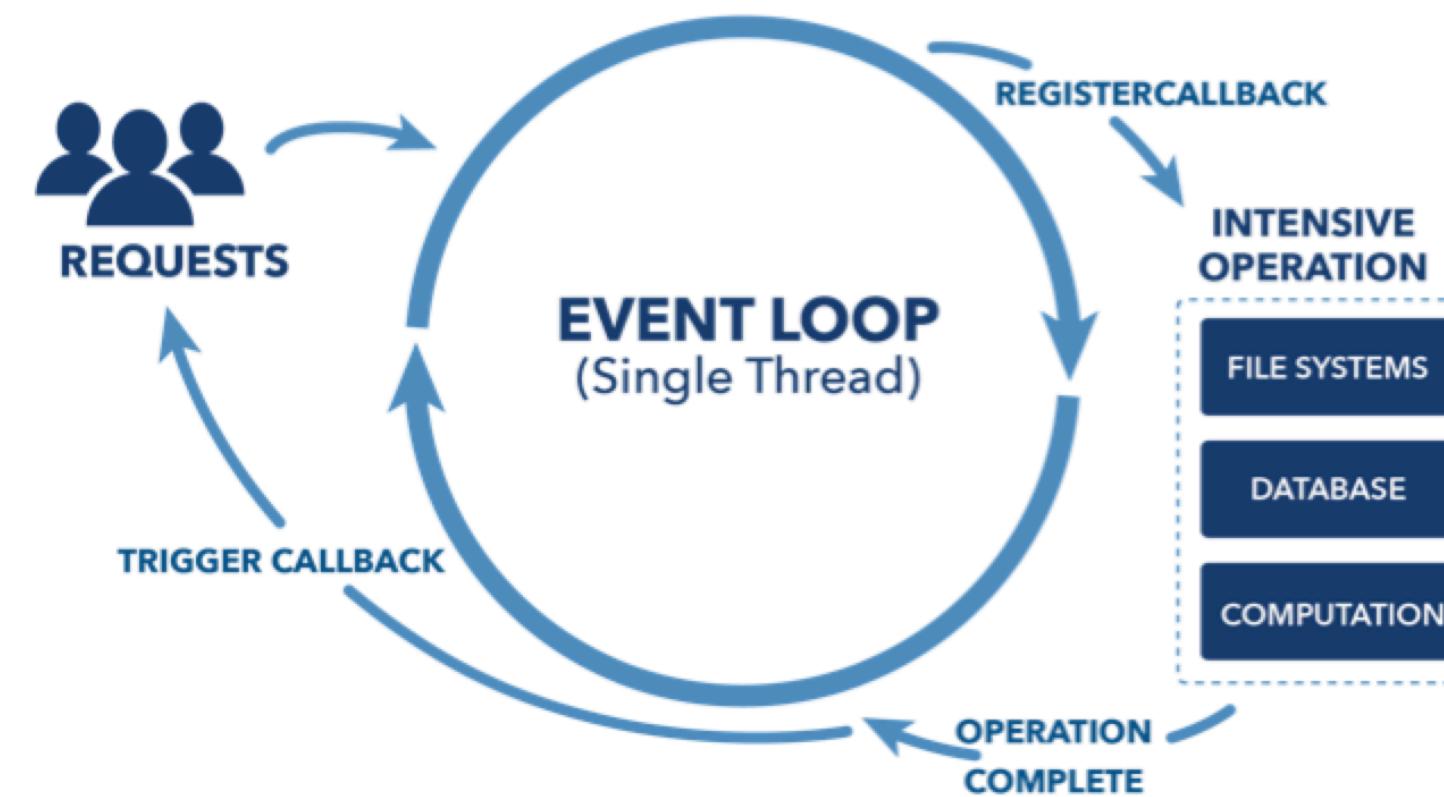
Input/output is one of
the most expensive
type tasks (>CPU)



Node has non-
blocking I/O

 TECH |

NON-BLOCKING I/O



Java Sleep

```
System.out.println("Step: 1");
System.out.println("Step: 2");
Thread.sleep(1000);
System.out.println("Step: 3");
```

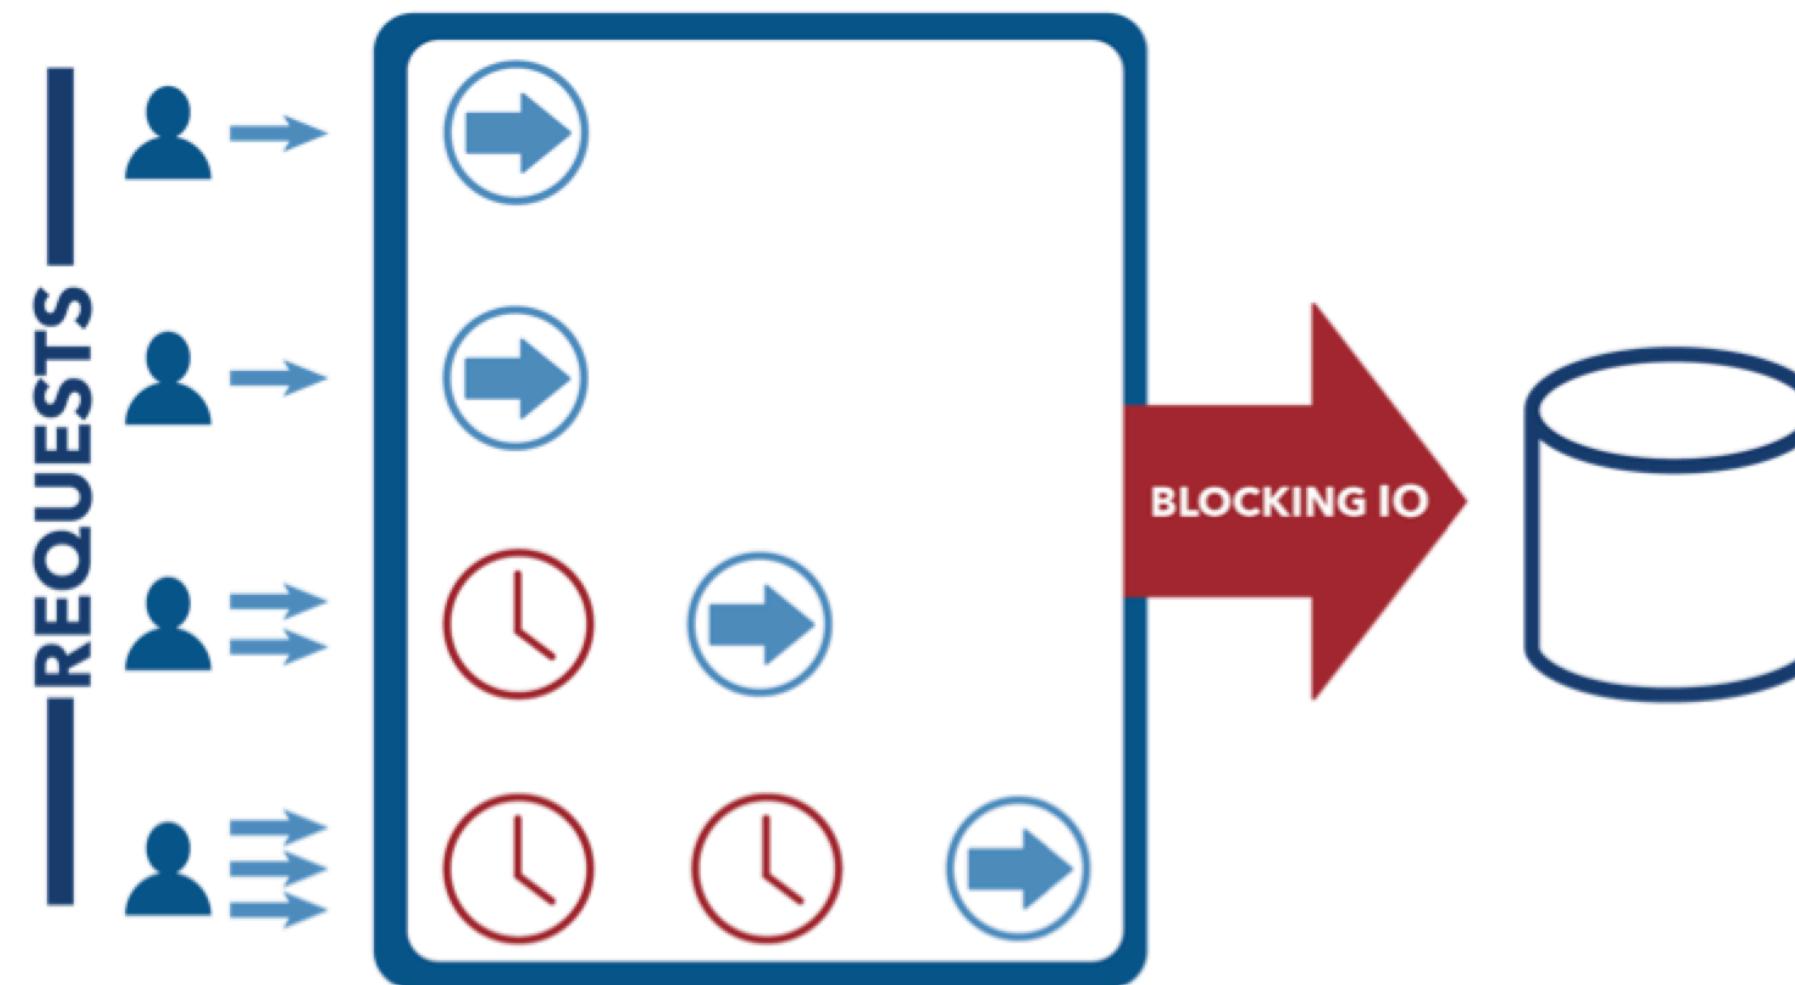
Node "Sleep"

```
console.log('Step: 1')
setTimeout(function () {
  console.log('Step: 3')
}, 1000)
console.log('Step: 2')
```

Process Multiple Tasks

```
console.log('Step: 1')
setTimeout(function () {
  console.log('Step: 3')
  // console.log('Step 5')
}, 1000);
console.log('Step: 2')
// console.log('Step 4')
```

Blocking Web Server

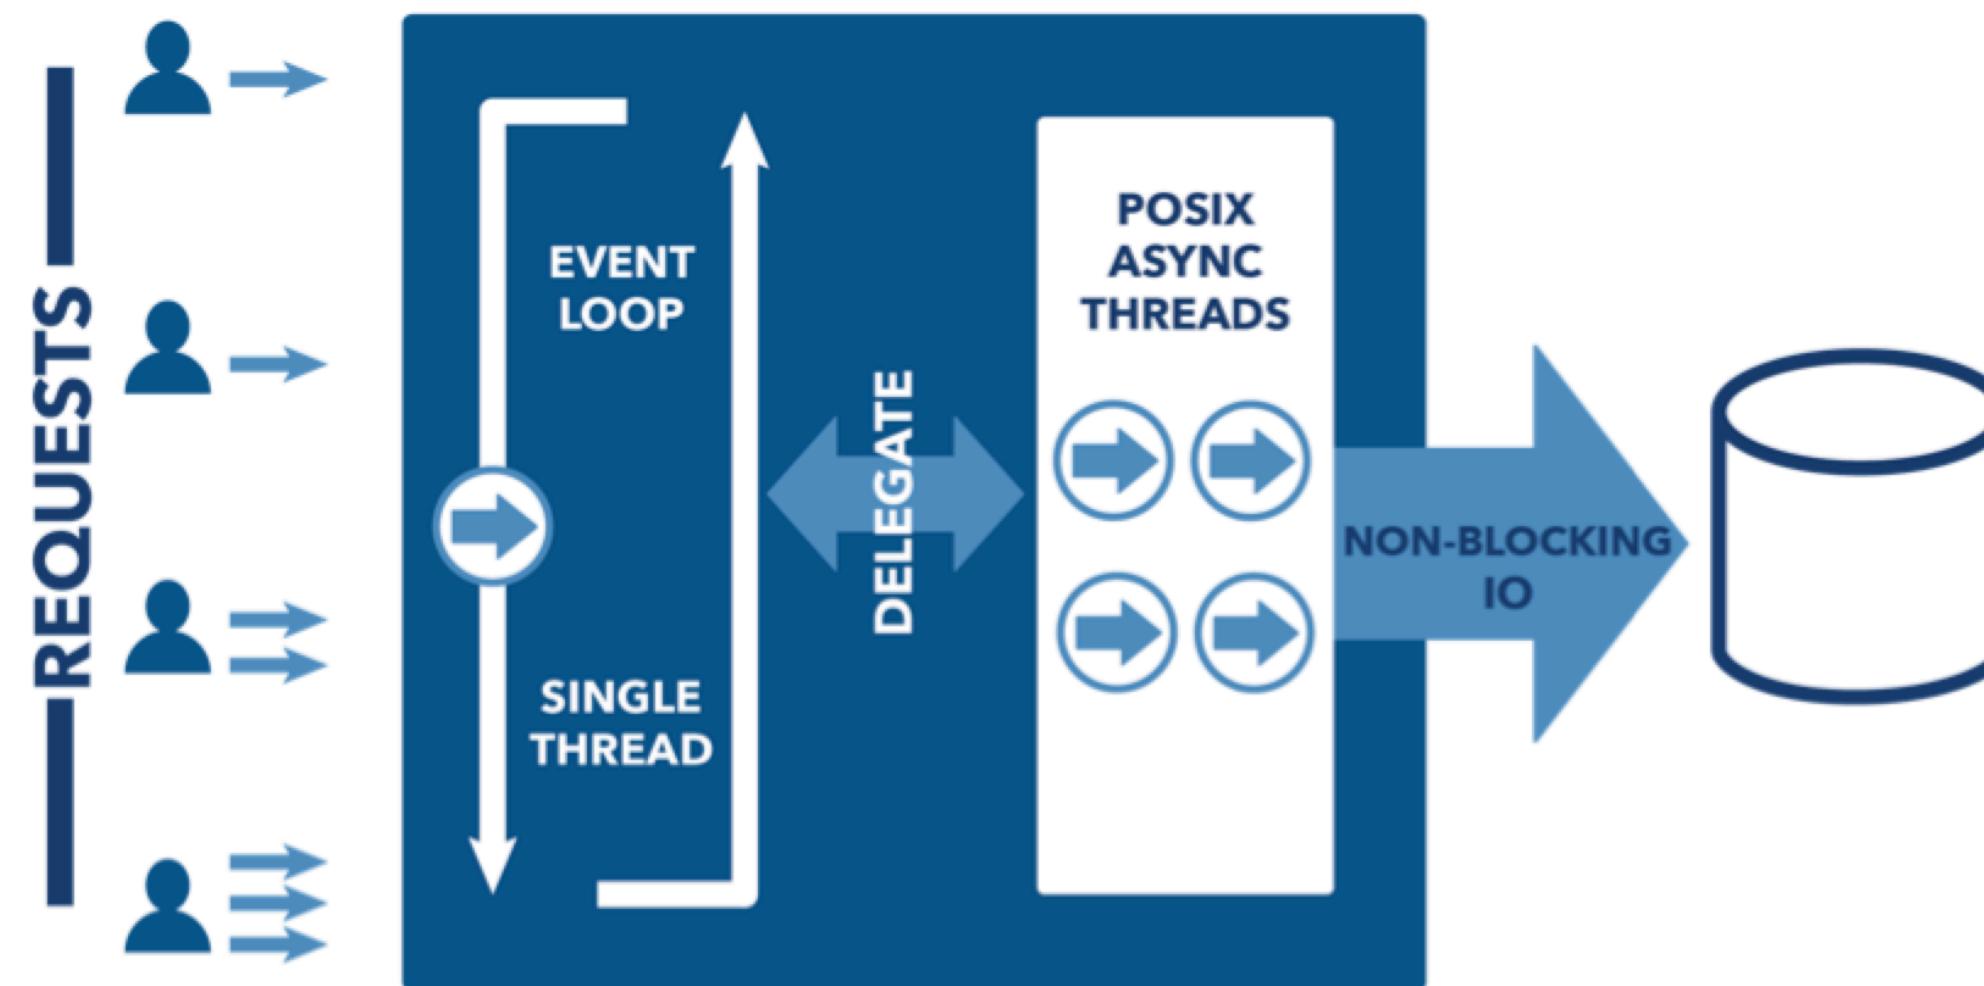
 TECH | **BLOCKING I/O**



Non-Blocking Web Server



TECH | NON-BLOCKING I/O





[Multi-threading] is the software equivalent of a nuclear device because if it is used incorrectly, it can blow up in your face.

<http://blog.codinghorror.com/threading-concurrency-and-the-most-powerful-psychokinetic-explosive-in-the-univ>

Blocking systems have
to be multi-threaded

Node is single
threaded... and that's
good! 😊

It's still possible to write blocking
code in Node.js. 😳

Blocking Node.js Code

```
// blocking.js  
console.log('Step: 1')  
for (var i = 1; i<1000000000; i++) {  
    // This will take 100-1000ms  
}  
console.log('Step: 2')
```

Blocking Node.js Code

```
var fs = require('fs')

var contents = fs.readFileSync('accounts.txt','utf8')
console.log(contents)
console.log('Hello Ruby\n')

var contents = fs.readFileSync('ips.txt','utf8')
console.log(contents)
console.log('Hello Node!')

//accounts.txt->Hello Ruby->ips.txt->Hello Node!
```

Non-Blocking Node.js Code

```
var fs = require('fs')

fs.readFile('accounts.txt','utf8', function(error, contents){
  console.log(contents)
})

console.log('Hello Ruby\n')

fs.readFile('ips.txt','utf8', function(error, contents){
  console.log(contents)
})

console.log('Hello Node!')

//Hello Ruby->Hello Node->... accounts.txt->ips.txt or ips.txt->accounts.txt
```

Node *typically* is much
faster than other
platforms

How many of you
reach the performance
limitations of apps
built with blocking I/O
systems?

Probably not many

My Fav Node Benefit

JavaScript everywhere.
One language to rule
'em all!

- >> Think faster
- >> Reuse code
- >> Learn quicker

Most of Node is JavaScript

- >> Array
- >> String
- >> Primitives
- >> Functions
- >> Objects

Node !== Browser
JavaScript

How to create global variables (no window in Node), work with modules, get path to my script?

global or GLOBAL

It has properties!

global.__filename
global.__dirname

```
global.module  
global.require()
```

How do I...?

- >> Access CLI input?
- >> Get system info: OS, platform, memory usage, versions, etc.?
- >> Read env vars (passwords!)?

global.processor
processor

process.pid
process.versions
process.arch

process.argv

process.env

```
process.uptime()  
process.memoryUsage()
```

`process.cwd()`

```
process.exit()  
process.kill()
```

Who likes and
understands
callbacks? 

http://callbackhell.com

```
fs.readdir(source, function (err, files) {
  if (err) {
    console.log('Error finding files: ' + err)
  } else {
    files.forEach(function (filename, fileIndex) {
      console.log(filename)
      gm(source + filename).size(function (err, values) {
        if (err) {
          console.log('Error identifying file size: ' + err)
        } else {
          console.log(filename + ' : ' + values)
          aspect = (values.width / values.height)
          widths.forEach(function (width, widthIndex) {
            height = Math.round(width / aspect)
            console.log('resizing ' + filename + 'to ' + height + 'x' + height)
            this.resize(width, height).write(dest + 'w' + width + '_' + filename, function(err) {
              if (err) console.log('Error writing file: ' + err)
            })
          }).bind(this))
        }
      })
    })
  }
})
```

Callbacks are not very
developmental scalable 😞

Me When Working With Deeply Nested Callbacks



Events

Events are part of core and supported by most of the core modules while more advanced patterns such as promises, generators, `async/await` are not.

Events == Node Observer Pattern

- >> Subject
- >> Observers (event listeners) on a subject
- >> Event triggers

Events

```
var events = require('events')
var emitter = new events.EventEmitter()
```

Events

In node.js an event can be described simply as a string with a corresponding callback.

```
emitter.on('done', function(results) {  
  console.log('Done: ', results)  
})
```

Using Event Emitters

```
var events = require('events')
var emitter = new events.EventEmitter()

emitter.on('knock', function() {
  console.log('Who\'s there?')
})

emitter.on('knock', function() {
  console.log('Go away!')
})

emitter.emit('knock')
```

Inheriting from EventEmitter

```
// job.js
var util = require('util')
var Job = function Job() {
  // ...
  this.process = function() {
    // ...
    job.emit('done', { completedOn: new Date() })
  }
}

util.inherits(Job, require('events').EventEmitter)
module.exports = Job
```

Inheriting from EventEmitter

```
// weekly.js  
  
var Job = require('./job.js')  
  
var job = new Job()  
  
job.on('done', function(details){  
  console.log('Job was completed at', details.completedOn)  
  job.removeAllListeners()  
})  
  
job.process()
```

Listeners

`emitter.listeners(eventName)`

`emitter.on(eventName, listener)`

`emitter.once(eventName, listener)`

`emitter.removeListener(eventName, listener)`

Other Node Patterns

Node Patterns: From Callbacks to Observer: [http://
webapplog.com/node-patterns-from-callbacks-to-observer/](http://webapplog.com/node-patterns-from-callbacks-to-observer/)

or

<https://github.com/azat-co/node-patterns>

Problems with Large Data

- » Speed: Too slow because has to load all
- » Buffer limit: ~1Gb
- » Overhyped (JK)

Streams

Abstractions for continuous
chunking of data

No need to wait for the
entire resource to load

Types of Streams

- » Readable
- » Writable
- » Duplex
- » Transform

Streams Inherit from Event Emitter

Streams are Everywhere!

- » HTTP requests and responses
- » Standard input/output (stdin&stdout)
- » File reads and writes

Readable Stream Example

`process.stdin`

Standard input streams contain data going into applications.

This is achieved via a
read operation.

Input typically comes from the keyboard used to start the process.

To listen in on data from stdin, use the data and end events:

```
// stdin.js

process.stdin.resume()
process.stdin.setEncoding('utf8')

process.stdin.on('data', function (chunk) {
  console.log('chunk: ', chunk)
})

process.stdin.on('end', function () {
  console.log('--- END ---')
})
```

Demo

```
$ node stdin.js
```

New Interface `read()`

```
var readable = getReadableStreamSomehow()

readable.on('readable', () => {
  var chunk

  while (null !== (chunk = readable.read())) {
    console.log('got %d bytes of data', chunk.length)
  }
})
```

Writable Stream Example

`process.stdout`

Standard output streams contain data going out of the applications.

This is done via a write operation.

Data written to
standard output is
visible on the
command line.

Writable Stream

To write to stdout, use the `write` function:

```
process.stdout.write('A simple message\n')
```

What about HTTP?

```
const http = require('http')

var server = http.createServer( (req, res) => {
    req.setEncoding('utf8')
    req.on('data', (chunk) => {
        transform(chunk) // This functions is defined somewhere else
    })
    req.on('end', () => {
        var data = JSON.parse(body)
        res.end()
    })
})

server.listen(1337)
```

Pipe

```
var r = fs.createReadStream('file.txt')
var z = zlib.createGzip()
var w = fs.createWriteStream('file.txt.gz')
r.pipe(z).pipe(w)
```

What data type to use for binary data?

Buffers

Binary data type, to create:

```
>> Buffer.alloc(size)
```

```
>> Buffer.from(array)
```

```
>> Buffer.from(buffer)
```

```
>> Buffer.from(str[, encoding])
```

Docs: <http://bit.ly/1IeAcZ1>

Working with Buffer

```
// buf.js
var buf = Buffer.alloc(26)
for (var i = 0 ; i < 26 ; i++) {
  buf[i] = i + 97 // 97 is ASCII a
}
console.log(buf) // <Buffer 61 62 63 64 65 66 67 68 69 6a 6b 6c 6d 6e 6f 70 71 72 73 74 75 76 77 78 79 7a>
console.log(buf.toString('utf8')) // abcdefghijklmnopqrstuvwxyz
```

Buffer Conversion

```
buf.toString('ascii') // outputs: abcdefghijklmnopqrstuvwxyz
```

```
buf.toString('ascii', 0, 5) // outputs: abcde
```

```
buf.toString('utf8', 0, 5) // outputs: abcde
```

```
buf.toString(undefined, 0, 5) // encoding defaults to 'utf8', outputs abcde
```

Remember fs?

```
fs.readFile('/etc/passwd', function (err, data) {  
  if (err) return console.error(err)  
  console.log(data)  
});
```

data is buffer!

Demo

```
$ node server-stream
```

Streams and Buffer Demo

```
// server-stream.js  
  
app.get('/stream', function(req, res) {  
  var stream = fs.createReadStream(largeImagePath)  
  stream.pipe(res)  
})
```

```
$ node server-stream
```

<http://localhost:3000/stream>

<http://localhost:3000/non-stream>

Results in DevTools

/stream responds faster!

X-Response-Time

~300ms vs. 3-5s

Stream Resources

Stream automated workshop: [https://github.com/substack/
stream-adventure](https://github.com/substack/stream-adventure)

```
$ sudo npm install -g stream-adventure  
$ stream-adventure
```

<https://github.com/substack/stream-handbook>

How to scale a single threaded system?

Cluster Usage

- >> Master: starts workers
- >> Worker: do the job, e.g., HTTP server

Number of processes = number of CPUs

Clusters

```
var cluster = require('cluster')

if (cluster.isMaster) {
  for (var i = 0; i < numCPUs; i++) {
    cluster.fork()
  }
} else if (cluster.isWorker) {
  // your server code
})
```

Cluster Demo

1. Run code/cluster.js with node (\$ node cluster.js).
2. Install loadtest with npm: \$ npm install -g loadtest
3. Run load testing with: \$ loadtest http://localhost:3000 -t 20 -c 10

Press control+c on the server terminal

Cluster Libraries

- » Core cluster: lean and mean
- » strong-cluster-control (<https://github.com/strongloop/strong-cluster-control>), or `\\$ slc run`: good choice
- » pm2 (<https://github.com/Unitech/pm2>): good choice

pm2

<https://github.com/Unitech/pm2>

<http://pm2.keymetrics.io>

Advantages:

- » Load-balancer and other features
- » os reload down-time, i.e., forever alive
- » Good test coverage

pm2 Demo: Typical Express Server

```
var express = require('express')
var port = 3000
global.stats = {}
console.log('worker (%s) is now listening to http://localhost:%s',
  process.pid, port)
var app = express()
app.get('*', function(req, res) {
  if (!global.stats[process.pid]) global.stats[process.pid] = 1
  else global.stats[process.pid] += 1;
  var l ='cluser '
    + process.pid
    + ' responded \n';
  console.log(l, global.stats)
  res.status(200).send(l)
})
app.listen(port)
```

pm2 Demo

Using server.js:

```
$ pm2 start server.js -i 0
```

In a new window:

```
$ loadtest http://localhost:3000 -t 20 -c 10  
$ pm2 list
```

Spawn vs Fork vs Exec

- » `require('child_process').spawn()` - large data, stream, no new V8 instance
- » `require('child_process').fork()` - new V8 instance, multiple workers
- » `require('child_process').exec()` - buffer, async, all the data at once

Spawn Example

```
fs = require('fs')
process = require('child_process')
var p = process.spawn('node', 'program.js')
p.stdout.on('data', function(data)) {
  console.log('stdout: ' + data)
})
```

Fork Example

```
fs = require('fs')
process = require('child_process')
var p = process.fork('program.js')
p.stdout.on('data', function(data)) {
  console.log('stdout: ' + data)
})
```

Exec Example

```
fs = require('fs')
process = require('child_process')
var p = process.exec('node program.js', function (error, stdout, stderr) {
  if (error) console.log(error.code)
})
```

How to handle async errors?

Handling Async Errors

Event Loop: Async errors are harder to handle/debug, because system loses context of the error. Then, application crashes.

Try/catch is not good enough.

Synchronous Error in Node

```
try {
    throw new Error('Fail!')
} catch (e) {
    console.log('Custom Error: ' + e.message)
}
```

For sync errors try/catch works fine.

Async Error Example

```
try {  
  setTimeout(function () {  
    throw new Error('Fail!')  
  }, Math.round(Math.random()*100))  
} catch (e) {  
  console.log('Custom Error: ' + e.message)  
}
```

The app crashes!

Me When Async Error's Thrown



Async Errors

How to deal with it?



Best Practices for Async Errors?

- » Listen to all “on error” events
- » Listen to uncaughtException
- » Use domain (soft deprecated) or AsyncWrap
- » Log, log, log & Trace
- » Notify (optional)
- » Exit & Restart the process

```
on('error')
```

Anything that inherits from or creates an instance of the above:
Express, LoopBack, Sails, Hapi, etc.

```
server.on('error', function (err) {  
  console.error(err)  
})
```

on('error') Chained Method Example

```
var http = require('http')
var server = http.createServer(app)
.on('error', function(e) {
  console.log('Failed to create server')
  console.error(e)
  process.exit(1)
})
```

on('error') Named Variable Example

```
var req = http.request(options, function(res) {  
    // ... processing the response  
})  
  
req.on('error', function(e) {  
    console.log('problem with request: ' + e.message)  
})
```

uncaughtException

uncaughtException is a very crude mechanism for exception handling. An unhandled exception means your application - and by extension Node.js itself - is in an undefined state. Blindly resuming means anything could happen.

uncaughtException

Always listen to uncaughtException!

```
process.on('uncaughtException', handle)
```

or

```
process.addListener('uncaughtException', handle)
```

uncaughtException Expanded Examples

```
process.on('uncaughtException', function (err) {  
  console.error('uncaughtException: ', err.message)  
  console.error(err.stack)  
  process.exit(1)  
})
```

Or

```
process.addListener('uncaughtException', function (err) {  
  console.error('uncaughtException: ', err.message)  
  console.error(err.stack)  
  process.exit(1)
```

Domain

This module is softly deprecated in 4.0 (most likely will be separate from core module), but there's no alternatives in core as of now.

Domain Example

```
var domain = require('domain').create()  
domain.on('error', function(error){  
    console.log(error)  
})  
domain.run(function(){  
    throw new Error('Failed!')  
})
```

Domain with Async Error Demo

domain-async.js:

```
var d = require('domain').create()
d.on('error', function(e) {
  console.log('Custom Error: ' + e)
})
d.run(function() {
  setTimeout(function () {
    throw new Error('Failed!')
  }, Math.round(Math.random()*100))
});
```

C++ Addons

How to Write C/C++ binding for your IoT, hardware, drone, smartdevice, etc.?

Node and C++

Create the hello.cc file:

```
#include <node.h>

namespace demo {

using v8::FunctionCallbackInfo;
using v8::HandleScope;
using v8::Isolate;
using v8::Local;
using v8::Object;
using v8::String;
using v8::Value;
```

Node and C++

Create the hello.cc file:

```
void Method(const FunctionCallbackInfo<Value>& args) {
    Isolate* isolate = args.GetIsolate();
    args.GetReturnValue().Set(String::NewFromUtf8(isolate, "capital one"));
}

void init(Local<Object> exports) {
    NODE_SET_METHOD(exports, "hello", Method);
}

NODE_MODULE(addon, init)
}

// namespace demo
```

Creating binding.gyp

Create binding.gyp:

```
{  
  "targets": [  
    {  
      "target_name": "addon",  
      "sources": [ "hello.cc" ]  
    }  
  ]  
}
```

node-gyp

```
$ npm install -g node-gyp
```

<https://github.com/nodejs/node-gyp>

Configuring and Building

```
$ node-gyp configure
```

```
$ node-gyp build
```

Check for compiled .node files in build/Release/

C++ Addons Examples

<https://github.com/nodejs/node-addon-examples>

Including Addon

Create hello.js and include your C++ addon:

```
var addon = require('./build/Release/addon')  
console.log(addon.hello()) // 'capital one'
```

Run

```
$ node hello.js
```

Want to work with Node but your boss won't let you?

Capital One is hiring ~2,000 more software engineers in UK,
Canada and US.

<https://jobs.capitalone.com>

We use Node and other cutting-edge open source tech a lot!
(React, Kotlin, Clojure, Angular 2, TypeScript, Go, etc.)

Learn More

Node at Capital One by Azat Mardan at Node Interactive 2015

<https://www.youtube.com/watch?v=BJPeLJhv1Ic>



30-Second Summary

1. Event Emitters
2. Streams
3. Buffers
4. Clusters
5. C++ Addons
6. Domain

Slides & Code



Everything: <https://github.com/azat-co/you-dont-know-node>

or

just PDF: <http://bit.ly/1VJWpQK>

My Contacts

Twitter: @azat_co

Email: hi@azat.co

Want to learn more about Node.js?

Check out Webapplog.com and NodeProgram.com for the best
online and in-person education!

One Last Thing

CodingHorror.com



Atwood's Law

Any application that *can* be written in JavaScript,
will eventually be written in JavaScript.